

## CS6650 Assignment 3 Report

Shengdi Wang - wang.shengd@northeastern.edu

### 1. GitHub Repo URL

<https://github.com/leowang396/distributed-music-service>

### 2. A 1-2 page description of your server design. Include major classes, packages, relationships, how messages get sent/received, etc.

The same sets of technologies are used for the web server and database layer, namely Tomcat/Java Servlets and Amazon DynamoDB. Since we need to incorporate asynchronous processing for likes and dislikes, I added a message broker server (running the RabbitMQ server) and a worker server for processing the message queue in this assignment.

The main packages used are RabbitMQ Java Client, AWS SDK for DynamoDB, Gson, and Java Servlet.

#### Web Server/API Layer

The server handles HTTP requests from the client. Calls to the albums API are processed synchronously as in Assignment 2. On the other hand, calls to the review API are processed asynchronously by publishing messages to a RabbitMQ queue. This decouples the processing from the main request-response cycle and allows the server to respond quickly to the client.

To handle the new review API on the web server, I added a new servlet (*ReviewServlet* class) to the Tomcat server. To manage and improve the efficiency of server resource usage, I also added the *RMQChannelFactory* class which is a RabbitMQ channel factory.

#### Message Broker Server

RabbitMQ is used as the message broker. By running it on a separate server, I ensure that the RMQ server has enough file descriptors available for establishing potentially many connections with both the Web Server and the Worker Server. The message broker server receives messages published by the web server and pushes them to available consumer channels on the same queue.

#### Worker Server for Message Queue

The worker server consumes messages from the same queue that the web server publishes to. Whenever a message is pushed to a worker thread by the message broker, it processes the data by calling *updateItem* API in DynamoDB. In this server design, the results of the worker operations are not communicated with the client.

There are 2 major classes for the worker server. The *Main* class establishes multiple channels with the message broker using multiple threads, ensuring that there is only 1 channel per thread and every message is properly acknowledged or rejected. Since the worker server

interfaces with the DynamoDB directly, I also added a *mqConsumerDdbDao* class which handles DynamoDB connection and requests.

#### Database

Similar to Assignment 2, I used DynamoDB as the database layer with pre-configured and pre-tuned read/write capacities, which ensures that there is no delay due to the auto-scaling of DynamoDB.

album
album_id
artist
title
year
image
likes
dislikes

*Figure 1: Assignment3 DynamoDB Data Model*

The data model is updated to include 2 new columns, “likes” and “dislikes”, which count the number of like and dislike API calls made against the albums. By storing the columns in the same table as other attributes, the review API operations can be fulfilled in a single DynamoDB API call, instead of needing a read call followed by a write call to check whether the album ID is valid before writing it.

3. Test run results (command lines showing metrics, RMQ management windows showing queue size, send/receive rates) showing your best throughput.

My best throughput for each configuration is summarized in the table below. Under all configurations, the current server design achieves a throughput of > 1000 requests per second, which is comparable to that of Assignment 2.

Configuration	Wall Time (s)	Throughput
threadGroupSize = 10, numThreadGroups = 10, delay = 2	32	1250
threadGroupSize = 10, numThreadGroups = 20, delay = 2	53	1509
threadGroupSize = 10, numThreadGroups = 30, delay = 2	85	1411

Below are the screenshots for the command lines and RMQ management windows, which demonstrate that the RMQ queue lengths are kept below 25 throughout all test runs. This indicates that with 50 Tomcat server channels and 200 RabbitMQ consumer channels, the messages are processed efficiently, and the message queue remains small to be stable.

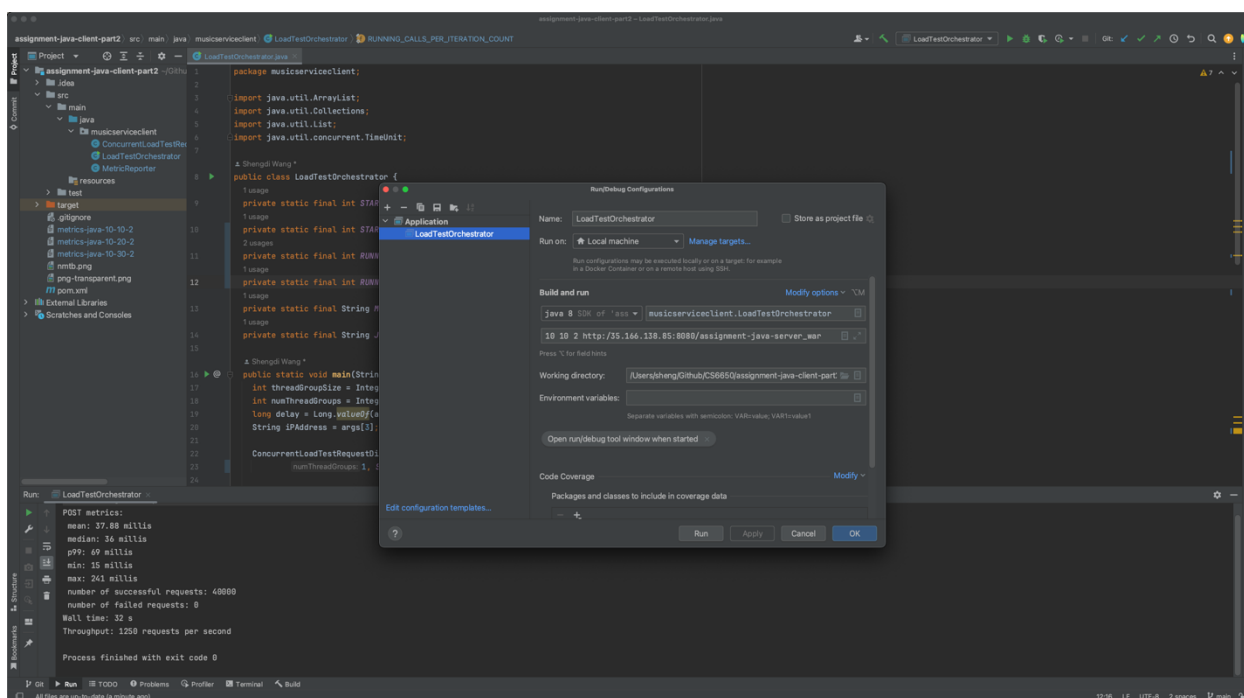


Figure 2: Command line screenshot for numThreadGroups = 10

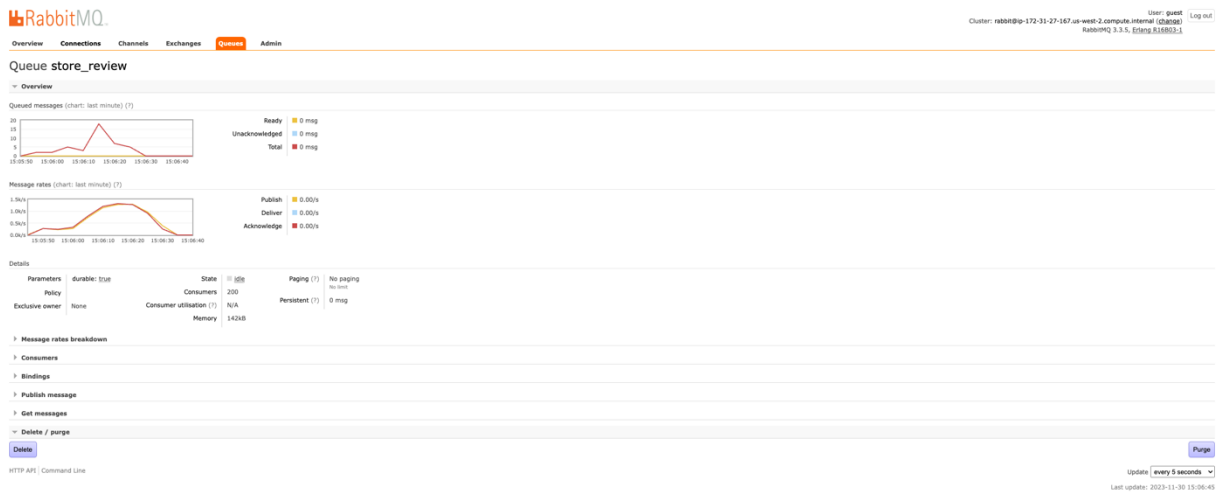


Figure 3: RMQ screenshot for numThreadGroups = 10

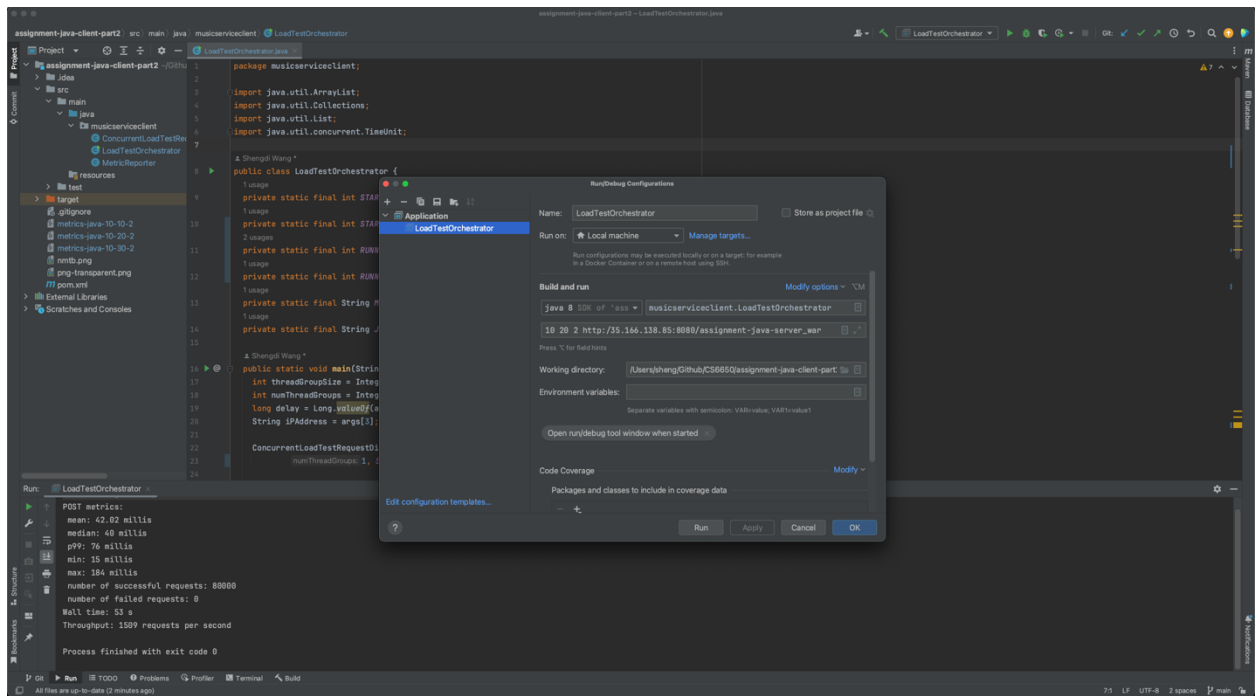


Figure 4: Command line screenshot for numThreadGroups = 20



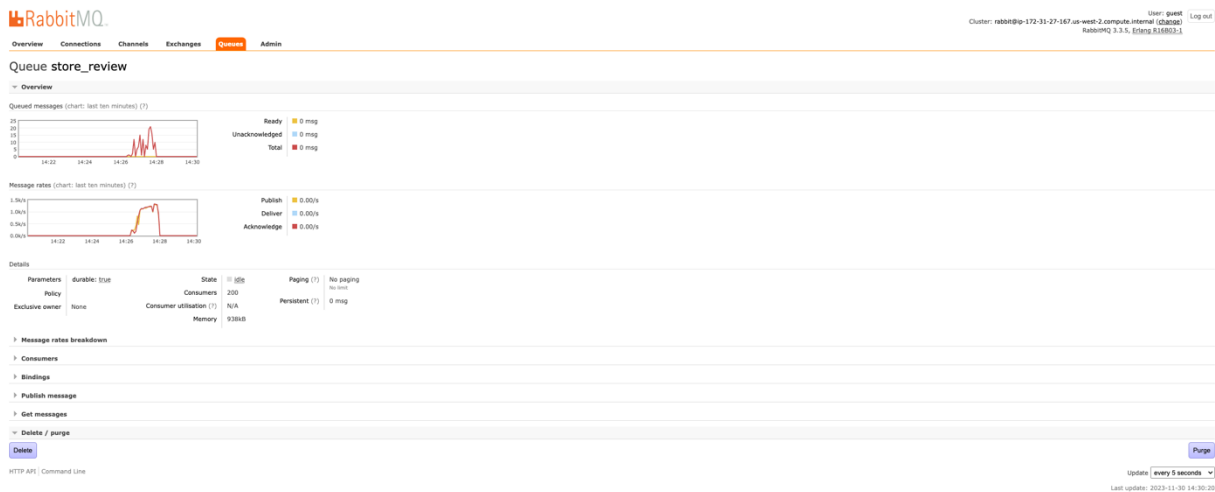


Figure 7: RMQ screenshot for numThreadGroups = 30